

# AI Service - Deployment Guide

Comprehensive guide for deploying the AI Service in various environments.

## Table of Contents

1. [Prerequisites](#)
2. [Local Development](#)
3. [Docker Deployment](#)
4. [Production Deployment](#)
5. [Kubernetes Deployment](#)
6. [Environment Configuration](#)
7. [Database Setup](#)
8. [Monitoring & Logging](#)
9. [Scaling](#)
10. [Troubleshooting](#)

## Prerequisites

### System Requirements

- **OS:** Linux (Ubuntu 20.04+ recommended) or macOS
- **CPU:** 2+ cores
- **RAM:** 4GB minimum, 8GB recommended
- **Storage:** 20GB minimum

### Software Requirements

- Python 3.11+
- PostgreSQL 14+
- Docker 20.10+ (for containerized deployment)
- Git

### External Services

- Abacus.AI account with API key
- Auth Service running and accessible
- Content Service (optional, for publishing)
- Social Media Service (optional, for social posting)

# Local Development

---

## 1. Clone Repository

```
git clone https://github.com/mission-engadi/ai-service.git
cd ai-service
```

## 2. Create Virtual Environment

```
python -m venv venv
source venv/bin/activate # On Windows: venv\Scripts\activate
```

## 3. Install Dependencies

```
pip install -r requirements.txt
pip install -r requirements-dev.txt # For development tools
```

## 4. Configure Environment

```
cp .env.example .env
# Edit .env with your configuration
```

### Required Variables:

```
DATABASE_URL=postgresql+asyncpg://user:password@localhost:5432/ai_service
SECRET_KEY=your-secret-key-here
JWT_SECRET_KEY=your-jwt-secret-key
ABACUS_API_KEY=your-abacus-api-key
```

## 5. Setup Database

```
# Create database
createdb ai_service

# Run migrations
alembic upgrade head
```

## 6. Start Service

```
# Using startup script
./scripts/start.sh

# Or manually
uvicorn app.main:app --host 0.0.0.0 --port 8010 --reload
```

## 7. Verify Installation

```
# Check status
./scripts/status.sh

# Test health endpoint
curl http://localhost:8010/api/v1/health

# Access API docs
open http://localhost:8010/docs
```

## Docker Deployment

### Using Docker Compose (Recommended)

#### 1. Configure environment:

```
cp .env.example .env
# Edit .env
```

#### 2. Start services:

```
docker-compose up -d
```

#### 3. View logs:

```
docker-compose logs -f ai-service
```

#### 4. Stop services:

```
docker-compose down
```

## Using Docker Directly

#### 1. Build image:

```
docker build -t mission-engadi/ai-service:latest .
```

#### 2. Run container:

```
docker run -d \
--name ai-service \
-p 8010:8010 \
-e DATABASE_URL="postgresql+asyncpg://user:pass@host:5432/ai_service" \
-e ABACUS_API_KEY="your-key" \
-e SECRET_KEY="your-secret" \
mission-engadi/ai-service:latest
```

#### 3. Check logs:

```
docker logs -f ai-service
```

## Production Deployment

### 1. Server Setup

**Update system:**

```
sudo apt update && sudo apt upgrade -y
```

**Install dependencies:**

```
sudo apt install -y python3.11 python3.11-venv postgresql nginx certbot
```

### 2. Application Setup

**Create application user:**

```
sudo useradd -m -s /bin/bash aiservice
sudo su - aiservice
```

**Clone and setup:**

```
git clone https://github.com/mission-engadi/ai-service.git
cd ai-service
python3.11 -m venv venv
source venv/bin/activate
pip install -r requirements.txt
```

**Configure environment:**

```
cp .env.example .env
vim .env
```

**Production .env:**

```
ENVIRONMENT=production
DEBUG=false
DATABASE_URL=postgresql+asyncpg://aiservice:password@localhost:5432/ai_service_prod
SECRET_KEY=<generate-strong-key>
JWT_SECRET_KEY=<generate-strong-key>
ABACUS_API_KEY=<your-api-key>
LOG_LEVEL=INFO
```

### 3. Database Setup

```
sudo -u postgres psql
```

```
CREATE DATABASE ai_service_prod;
CREATE USER aiservice WITH ENCRYPTED PASSWORD 'strong_password';
GRANT ALL PRIVILEGES ON DATABASE ai_service_prod TO aiservice;
\q
```

#### Run migrations:

```
alembic upgrade head
```

## 4. Systemd Service

#### Create service file:

```
sudo vim /etc/systemd/system/ai-service.service
```

```
[Unit]
Description=AI Service - Mission Engadi
After=network.target postgresql.service

[Service]
Type=simple
User=aiservice
Group=aiservice
WorkingDirectory=/home/aiservice/ai-service
Environment="PATH=/home/aiservice/ai-service/venv/bin"
EnvironmentFile=/home/aiservice/ai-service/.env
ExecStart=/home/aiservice/ai-service/venv/bin/uvicorn app.main:app --host 0.0.0.0 --port 8010 --workers 4
Restart=always
RestartSec=10

[Install]
WantedBy=multi-user.target
```

#### Enable and start:

```
sudo systemctl daemon-reload
sudo systemctl enable ai-service
sudo systemctl start ai-service
sudo systemctl status ai-service
```

## 5. Nginx Reverse Proxy

#### Create Nginx config:

```
sudo vim /etc/nginx/sites-available/ai-service
```

```

upstream ai_service {
    server 127.0.0.1:8010;
}

server {
    listen 80;
    server_name ai.mission-engadi.org;

    location / {
        proxy_pass http://ai_service;
        proxy_set_header Host $host;
        proxy_set_header X-Real-IP $remote_addr;
        proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
        proxy_set_header X-Forwarded-Proto $scheme;
        proxy_read_timeout 300;
        proxy_connect_timeout 300;
    }
}

```

#### **Enable site:**

```

sudo ln -s /etc/nginx/sites-available/ai-service /etc/nginx/sites-enabled/
sudo nginx -t
sudo systemctl reload nginx

```

## **6. SSL Certificate**

```
sudo certbot --nginx -d ai.mission-engadi.org
```

## **Kubernetes Deployment**

### **1. ConfigMap**

**configmap.yaml:**

```

apiVersion: v1
kind: ConfigMap
metadata:
  name: ai-service-config
  namespace: mission-engadi
data:
  ENVIRONMENT: "production"
  LOG_LEVEL: "INFO"
  SERVICE_PORT: "8010"

```

## 2. Secrets

```
kubectl create secret generic ai-service-secrets \
--from-literal=DATABASE_URL=postgresql+asyncpg://... \
--from-literal=SECRET_KEY=... \
--from-literal=JWT_SECRET_KEY=... \
--from-literal=ABACUS_API_KEY=... \
-n mission-engadi
```

## 3. Deployment

`deployment.yaml`:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: ai-service
  namespace: mission-engadi
spec:
  replicas: 3
  selector:
    matchLabels:
      app: ai-service
  template:
    metadata:
      labels:
        app: ai-service
    spec:
      containers:
        - name: ai-service
          image: mission-engadi/ai-service:latest
          ports:
            - containerPort: 8010
          envFrom:
            - configMapRef:
                name: ai-service-config
            - secretRef:
                name: ai-service-secrets
      resources:
        requests:
          memory: "512Mi"
          cpu: "250m"
        limits:
          memory: "2Gi"
          cpu: "1000m"
      livenessProbe:
        httpGet:
          path: /api/v1/health
          port: 8010
        initialDelaySeconds: 30
        periodSeconds: 10
      readinessProbe:
        httpGet:
          path: /api/v1/ready
          port: 8010
        initialDelaySeconds: 10
        periodSeconds: 5
```

## 4. Service

**service.yaml:**

```
apiVersion: v1
kind: Service
metadata:
  name: ai-service
  namespace: mission-engadi
spec:
  selector:
    app: ai-service
  ports:
    - port: 80
      targetPort: 8010
    type: ClusterIP
```

## 5. Ingress

**ingress.yaml:**

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: ai-service-ingress
  namespace: mission-engadi
  annotations:
    cert-manager.io/cluster-issuer: "letsencrypt-prod"
spec:
  tls:
    - hosts:
        - ai.mission-engadi.org
      secretName: ai-service-tls
  rules:
    - host: ai.mission-engadi.org
      http:
        paths:
          - path: /
            pathType: Prefix
        backend:
          service:
            name: ai-service
            port:
              number: 80
```

## 6. Deploy

```
kubectl apply -f k8s/
kubectl get pods -n mission-engadi
kubectl logs -f deployment/ai-service -n mission-engadi
```

# Environment Configuration

---

## Development

```
ENVIRONMENT=development
DEBUG=true
LOG_LEVEL=DEBUG
DATABASE_URL=postgresql+asyncpg://user:pass@localhost:5432/ai_service_dev
```

## Staging

```
ENVIRONMENT=staging
DEBUG=false
LOG_LEVEL=INFO
DATABASE_URL=postgresql+asyncpg://user:pass@staging-db:5432/ai_service_staging
```

## Production

```
ENVIRONMENT=production
DEBUG=false
LOG_LEVEL=WARNING
DATABASE_URL=postgresql+asyncpg://user:pass@prod-db:5432/ai_service_prod
```

---

# Monitoring & Logging

---

## Application Logs

### Using systemd:

```
sudo journalctl -u ai-service -f
```

### Using Docker:

```
docker logs -f ai-service
```

## Health Monitoring

```
# Health check
curl https://ai.mission-engadi.org/api/v1/health

# Readiness check
curl https://ai.mission-engadi.org/api/v1/ready
```

# Scaling

---

## Horizontal Scaling

**Docker Compose:**

```
docker-compose up -d --scale ai-service=3
```

**Kubernetes:**

```
kubectl scale deployment ai-service --replicas=5 -n mission-engadi
```

## Vertical Scaling

Update resources in deployment.yaml and apply.

---

# Troubleshooting

---

## Service Won't Start

**Check logs:**

```
sudo journalctl -u ai-service -n 100
```

**Common issues:**

- Database connection: Verify DATABASE\_URL
- Port in use: Check if port 8010 is available
- Missing dependencies: Reinstall requirements

## Database Connection Errors

```
# Test connection
psql $DATABASE_URL

# Check PostgreSQL status
sudo systemctl status postgresql
```

## High Memory Usage

- Reduce worker count
  - Enable database connection pooling
  - Implement request rate limiting
- 

**Last Updated:** December 2024

**Version:** 1.0.0